

Mythos Readiness: Why Discovery Was Never the Hard Part

Claude Mythos Preview solved vulnerability discovery. Scout is built for what happens next.

Scout Product Team · 2026

In April 2026, AWS, Apple, Google, Microsoft, and NVIDIA announced Project Glasswing — a defensive coalition built around Claude Mythos Preview, a frontier model purpose-built for deep code reasoning. Mythos doesn't just flag known CVEs faster than a scanner; it reads straight through opaque, complex codebases and surfaces high-severity vulnerabilities other tools never see at all.

That capability won't stay confined to a handful of frontier labs. Within the next 12–18 months, comparable models will run against ordinary enterprise pipelines — and the organizations that come out ahead won't be the ones with the best scanner. Discovery is already solved. Readiness is about what happens the moment a model like Mythos hands you hundreds of new, real findings at once: can you coordinate a response, close the blind spots it finds, filter the noise it repeats, and fix things at the speed it discovers them?

Four Scout capabilities carry that response end to end:

Campaigns

Closes the gap: mass CVE disclosure with no coordinated response path.

A 400-finding scan result shouldn't turn into a spreadsheet and a week of ownership disputes. Scout Campaigns auto-groups Mythos-scale findings by CVE, software, or asset, alerts notify groups and watchlists within the hour, and routes ownership through the same rules that already govern your findings — with velocity tracked live until the campaign closes.

BOM Grid

Closes the gap: fragmented bill-of-materials visibility across four BOM types.

Frontier models trace what your code depends on — including AI/ML models, cryptographic primitives, and vendor components most tooling never looks at. BOM Grid unifies SBOM, AI BOM, CBOM, and Vendor BOM into one exposure grid, so nothing Mythos flags lands in a blind spot.

Suppression Rules

Closes the gap: alert fatigue from repeat, previously-accepted findings.

A model that re-scans nightly will re-surface the same accepted-risk findings every run unless something filters them with a record of why. Scout keeps that filter governed instead of quiet: every suppressed finding carries a reason, an approver, and an expiry date, so yesterday's accepted risk doesn't quietly become tomorrow's blind spot.

Fix Intelligence

Closes the gap: remediation research lag behind AI-speed discovery.

Mythos can surface a critical vulnerability far faster than a human can research the fix. Fix Intelligence generates the concrete remediation — an upgrade, a patch, a config change — per CVE or software identity, prioritizes the queue by exploitability and SLA proximity, and returns structured fallback guidance instead of a dead end when AI confidence is low.

Under the Hood: CVE Investigation

All four capabilities lean on the same engine. The moment a CVE shows up against your inventory, Scout automatically runs false-positive checks, end-of-life analysis, and exploit-status lookups — before a human opens the finding.

That single step is what lets suppression decisions get made with confidence, feeds the recommendations behind Fix Intelligence, and turns a raw Mythos-flagged CVE into something a Campaign can assign as real work instead of noise. CVE Investigation is the foundation; the other four are what you build on top of it.

Mythos readiness isn't measured by how fast you can find what's wrong — that part is already solved. It's measured by how fast you can coordinate, close a blind spot, filter the repeat noise, and ship the fix. Scout is built to make that measurement come out in your favor.

Curious where your environment stands today? Talk to the Scout team about a readiness walkthrough.